

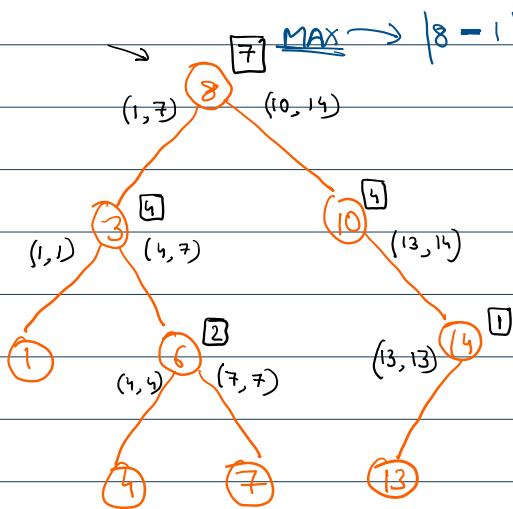
228.Summary ranges

12 June 2023 09:58 PM

1146.Snapshot array

12 June 2023 09:59 PM

Maximum diff. b/w node & ancestor :



$$\text{MAX} \rightarrow |8 - 1|$$

+ node,

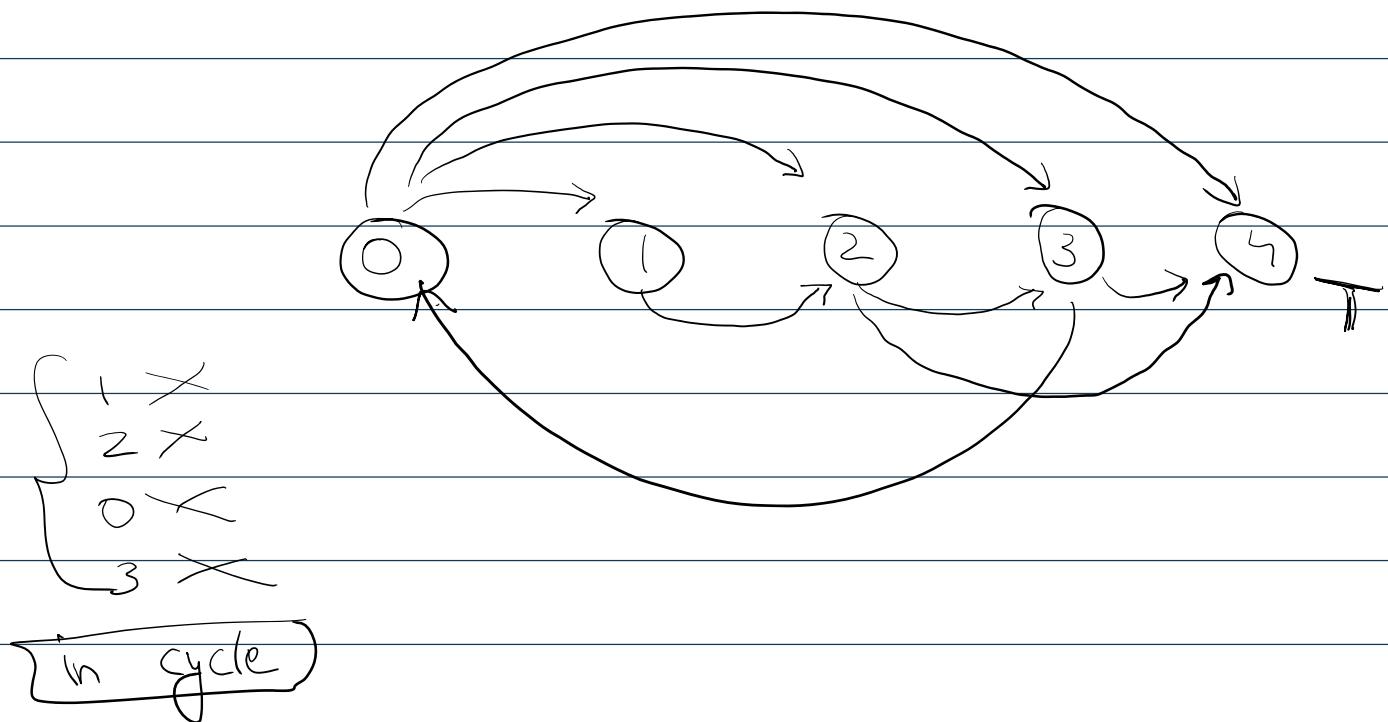
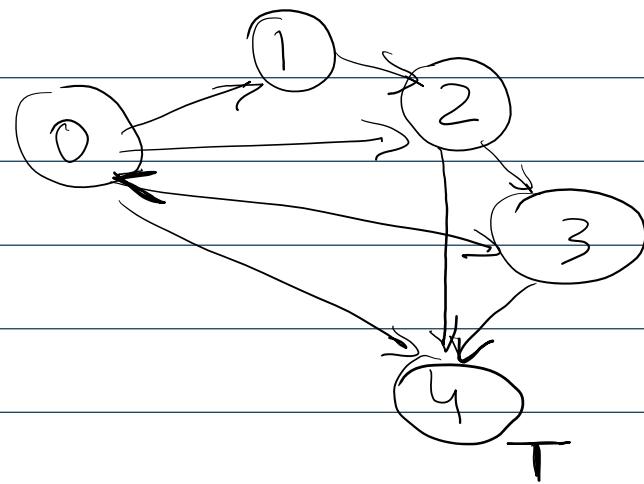
we'll find (min, max) from both its L_T & R_T .

Then compute V:

$$V = \left\{ \begin{array}{l} \min - \text{curr} \\ \max - \text{curr} \end{array} \right\} \text{Max.}$$

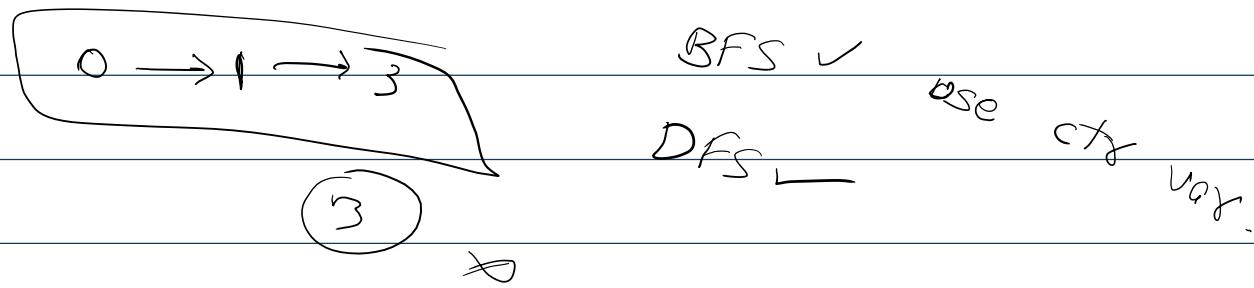
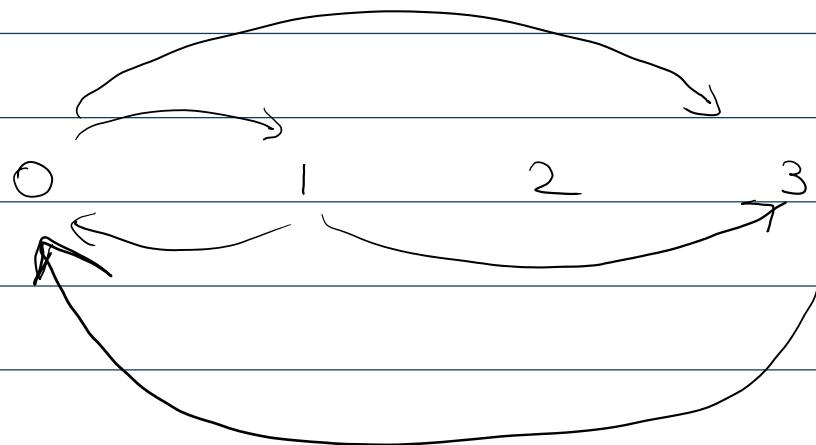
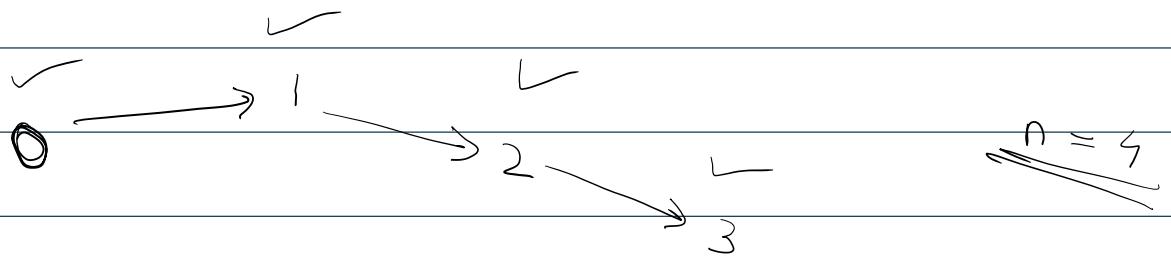
&

iterate till we find Max. (V)



841

4 September 2024 04:57 AM

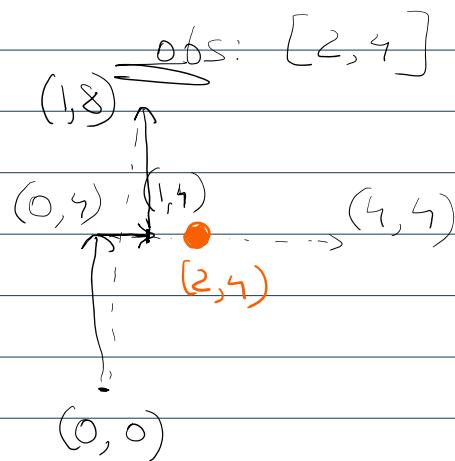
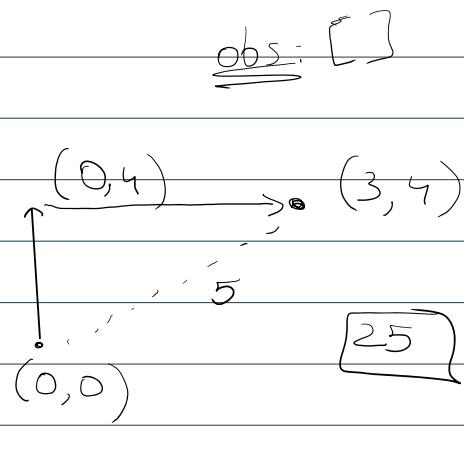


$$\begin{matrix} N \\ \downarrow \\ (0, 0) \end{matrix}$$

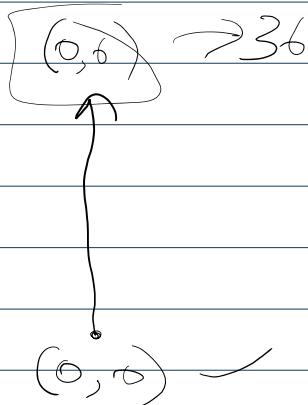
$$-2 : L, -1 : R$$

$$[1, 4] \in \mathbb{R}$$

$$\boxed{x^2 + y^2}$$



$$(\sqrt{1^2 + 8^2})^2 \rightarrow \boxed{65}$$





0 1 2 3
($[1, 3]$, $[3, 0, 1]$, $[2]$, $[0]$)

$$ct\gamma = \cancel{\alpha} \cancel{\beta}$$

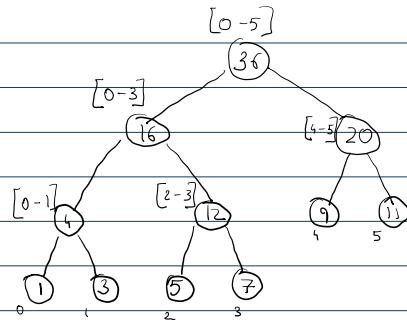
$$n = 4$$

vis	T	T	F	F
	0	1	2	3

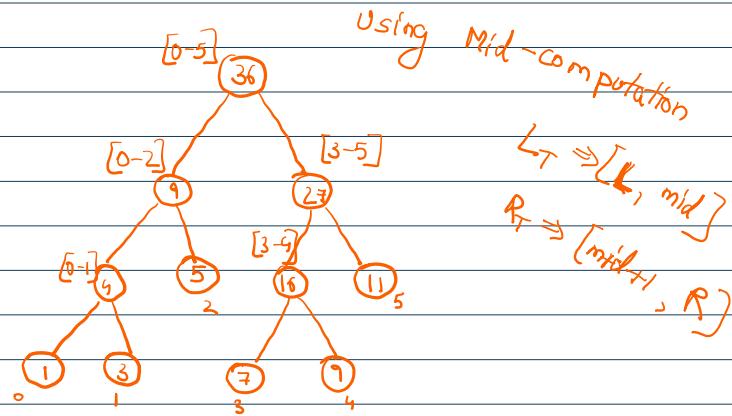
Segment tree

0 1 2 3 4 5
1, 3, 5, 7, 9, 11
1 4 9 16 25 36

$n = 6$



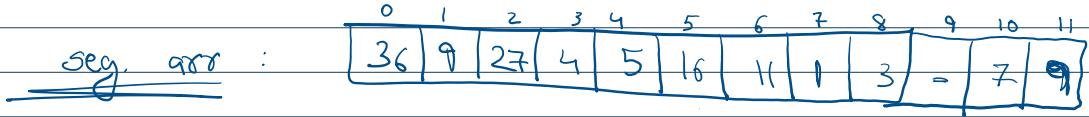
OR

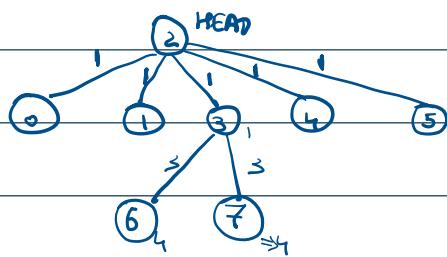


→ leaf nodes are arr. elements

$$\begin{aligned} L_c &= 2i + 1 \\ R_c &= 2i + 2 \end{aligned} \quad \left\{ \text{for Root@ } i \text{ index} \right.$$

nodes
 $2n - 1$





u v wt
mgozi's i information[u]

2 0 1

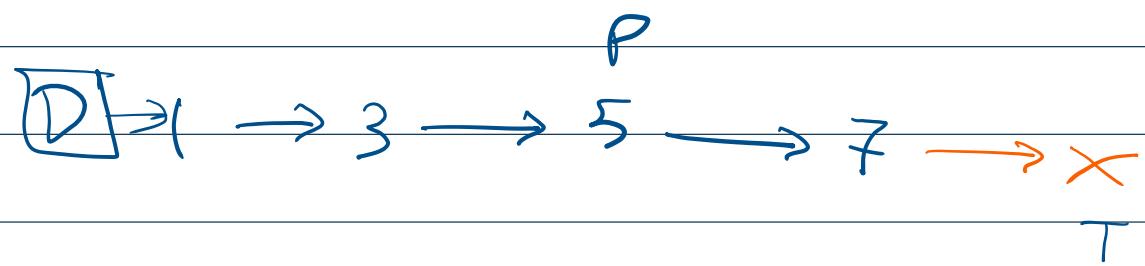
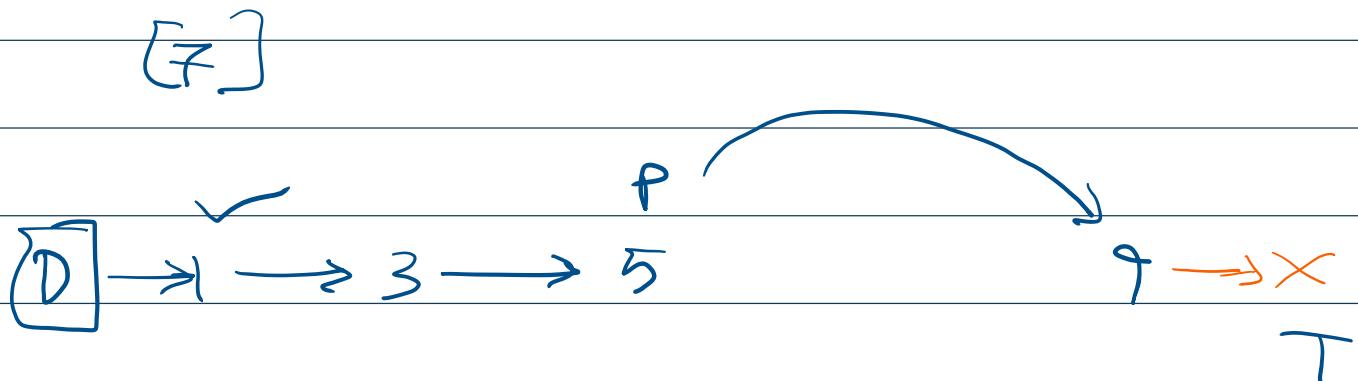
2 1 1

-1 2

information[com_mys]

3217

7 September 2024 12:27 AM



$$k = 4, n = 3$$

Max-heap (k)

$[1 \ 2 \ 3 \ 4 \ 5]$

diff.
w.r.t. n

2 1 0 1 2

$\{1, 2\}$

$\{5, 2\}$

$\{2, 1\}$

$\{3, 0\}$

$\{4, 1\}$

arr. $\rightarrow 1, 2, 3, 4, 5$, $n = 3$

$K = 4$

MIN HEAP

1 2 3 4 5

~~diff.~~ 2 1 0 1 2

{2, 1}

{2, 5}

$K = 4$

res : 3, 2, 4, 1

0 1 2 3 4

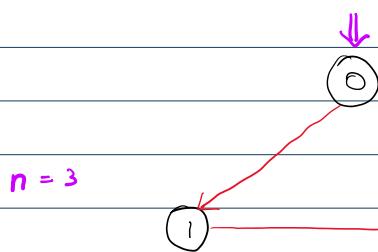
arr. $\rightarrow [1, 2, 3, 4], 5$

$K = 4$

~~diff.~~ 2 1 0 1 2

$n \approx 3$

↑ ↑ ↑
L R R



directed

$$n = \infty \rightarrow \text{ref } m$$

~~else~~

~~set (n, m% n)~~

$$\text{gcd}(18, 6)$$

$$(6, 0) \quad \boxed{6}$$

$$(10, 3)$$

$$(3, 1)$$

(1,0)

() ()

answer : $[0, 1, -1]$

$$(6, 10)$$

(10, 6)

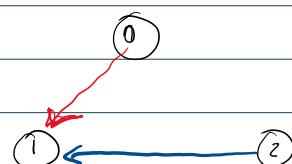
$$(6, 4)$$

$$(4, 2)$$

(2, 0)

1

$$\underline{n = 3}$$



$$\text{nodes} \quad 0 \quad 1 \quad 2$$

Moore's law

Transistors / chip will double roughly 1-2 yr.

Dennard Scaling → (power decipated)

→ Power density remains const.

→ Performance per watt ↑ exponentially.

SPEC → std. Performance Eval." Cooperative

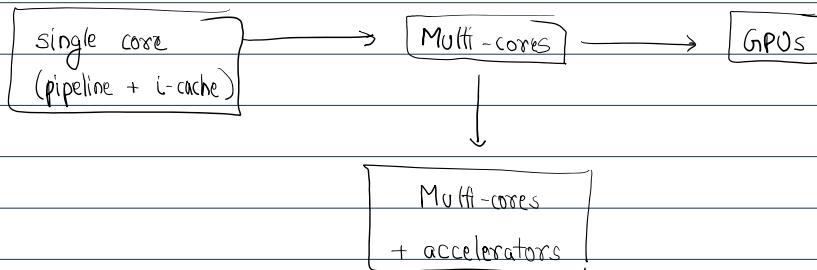
↑ in Performance:

gen'. IN-ORDER ← n-issue IN-ORDER ← n-issue OOO

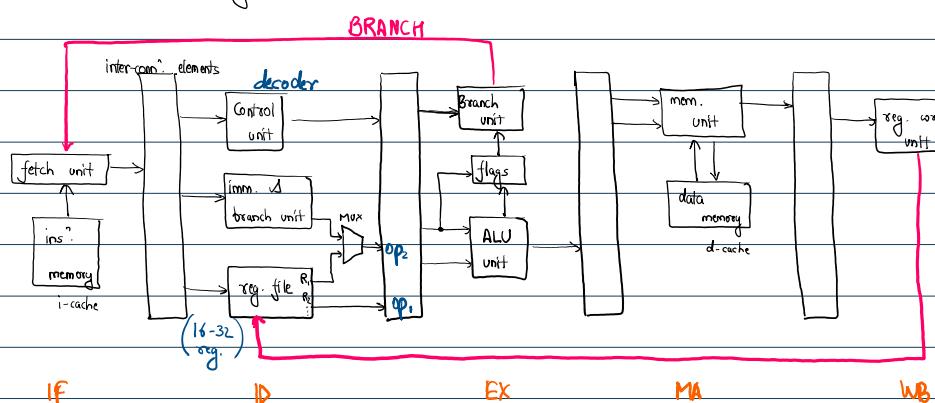
pipeline

pipeline

pipeline

Out-of-order Pipeline:

Processor with 5-stages:



Pipeline eliminates idleness in processor

{ except for time to fill & drain it. }

	<u>ins. load</u>	<u>exec.</u>	<u>commit</u>
In-order pipeline :	in-order	in-order	in-order
Out-of-order pipeline :	000	000	in-order

→ Pipelines use -ve edge triggered latches / buffer b/w the 2 stages, to store the imm. results & synchronize the clock.

* Hazards :

- i) Structural hazards → same resource / h/w access
- ii) Data hazards → data dependency (RAW, WAR, WAW)
- iii) Control hazards → Branch statements

Data interlocks :

identifies dependencies on ins. & inserts NOP ins. / bubbles whenever it finds suitable.
→ stalls the pipeline, Performance ↓

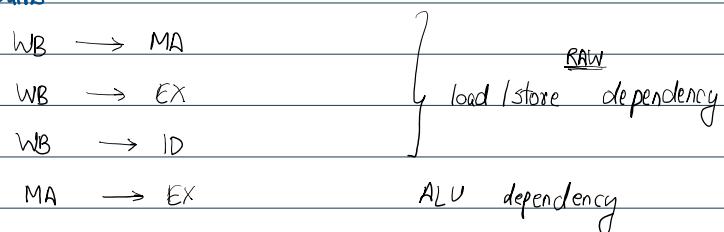
[Stalls]

→ Optimization for RAW dependency:

Forwarding : from MA to EX stage, resulting in NO stall,
{WB to reg. file, will happen in background}

forwarding unit uses MUX to determine, which i/p needs to be taken for ALU exec.

Types of forwarding paths :



* forwarding cannot help in case of LOAD-use hazard (RAW dependency)
↓
requires 1 stall cycle.

* dis-adv. of in-order pipelines :

- load-use hazard : causes 1 stall each
- Usually we have multicycle MA stage

- load-use hazard : causes 1 stall each
- usually we have multicycle MA stage
- Convey effect (Head of line) HOL blocking
- Taken branches have 2 stalls each
- IPC < 1

* Performance eq.

$$\frac{\# \text{Programs}}{\# \text{sec.}} = \frac{\text{IPC} \times \text{freq.}}{\# \text{ins}}, \text{ for 1 program}$$

or

$$\text{Performance} \propto \frac{\text{IPC} * \text{freq.}}{\# \text{ins}}$$

$$\text{IPC} = \frac{1}{\text{CPI}}$$

→ #ins depends on compiler.

→ freq. depends on transistor technology & arch. (#Pipeline stages)

(clock Period)

$$t_{clk} = \frac{t}{k} + \Delta, \text{ where, } k\text{-stage pipeline,}$$

with Δ as buffer delay.

$$f_{max} = \frac{1}{t + \Delta} \Rightarrow f_{max.} \propto k, \text{ give } \frac{t}{k} \gg \Delta$$

→ IPC again depends on arch. & compilers. { lower stalls, IPC ↑ }
 i.e., forwarding,
 rearranging code, keeping in mind data-dependency.

* Methods to ↑ IPC :

- forwarding
- having more not taken branches
- faster ins? & data memories

* ↑ freq. is LIMITED by

- i) Power, $P \propto f^3$
- ii) Change in temp, $\Delta T \propto P$
- iii) ↑ in pipeline stages \Rightarrow More hazards, more forwarding paths, design complexity
- iv) Effect of latch delay

* CPI_{actual} = CPI_{ideal} + (stall_rate * stall_penalty)

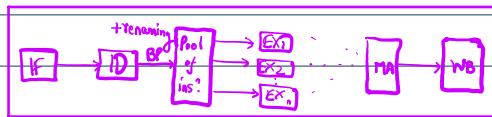
* Super-scalar Processor: helps ↑ IPC

→ can execute multiple insⁿ per cycle.
 { by issuing more insⁿ per cycle }
 i.e.

Intel Pentium : had 2 in-order pipelines (U & V)
 insⁿ : i_j i_{j+1}

dis-adv:

- dependency b/w insⁿ
- $O(n^2)$ forwarding paths for n-issue processor.
- complicated logic for detecting hazards, dependencies, & forwarding.



* Out-of-order Pipelines:

- don't follow program order for execⁿ
- { it is as per data dependence order }
- extracts more //ism. from code segment.
- { insⁿ lvl parallelism }

- Steps:
- i) Creates a pool of insⁿ { fetches: Program order }
 - ii) finds mutually independent insⁿ, having operands ready
 - iii) Executes them ooo

→ larger the POOL \Rightarrow more ILP can be exploited,
 insⁿ. window given insⁿ are on correct path
 (typical size: 16-128)

Branch Prediction:

Consider,

Answe,

$$\# \text{ins} \rightarrow n$$

$$\# \text{branches} \rightarrow n/5$$

$P(E)$ predicting any given branch INCORRECTLY $\rightarrow p$

Then,

$P(E)$ predicting ALL branches CORRECTLY $\rightarrow (1-p)^{n/5}$

$\Rightarrow P(E)$ of making atleast a single mistake $\rightarrow [1 - (1-p)^{n/5}]$
in a pool of 'n' ins".

\rightarrow We require Branch Prediction to be highly accurate.

i.e. it does have a very high substantial penalty.
 \Downarrow

even with 97.5% acc. in BP \Rightarrow we have 40% chances of
having incorrect ins". in window.

\rightarrow Acc. of BP limits the size of ins" window.

* Types of data-dependencies:

RAW \rightarrow true dependency { causes problem in-order pipeline }

(rp-dependency) WAW
(anti-dependency) WAR } false dependency, caused due to finite no. of reg.

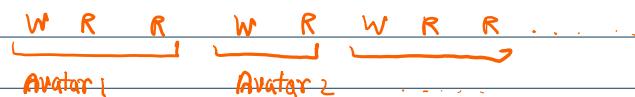
{ in 000 pipeline } Sol.: Register renaming

(done by HCo, internally)

Register renaming:

In the assembly eq^w code,
we map an architectural register to a physical reg.,
in-order to eliminate false dependencies.

>Create new avatars of reg. whenever we have a W op.
i.e,



(write once, read many times)

\rightarrow Only RAW dependencies exists in the code, (if any)
after renaming

→ Only RAW dependencies exists in the code, (if any)
after renaming.

↓
higher ILP

Precise exceptions :

→ we need precise behaviour of ooo pipeline,
to handle stop & resume effectively (CSW.), in case of an exception/
interrupt

Consider, intr. after k^{th} ins?

Then,

ins? till k should has executed completely

↓ written their results to MEM./Reg. file.

↓, $(k+1)^{\text{th}}$ ↓ later, should not appear
to have started their EX at all.

* IF :

To explore ILP (in ooo), we need to ↑ fetch BW.

↙ But it incurs other overhead, → we only know PC value
→ no time to look at ins? (prior decode in IF stage)

such as for [B.P] → It becomes difficult to predict correct
set of ins?, in a multi-way branch.

{nested branch structure }

for an optimal compiler, if we have a fetch BW = 4, we'll
expect it to place Branch statements as far as possible.

Branch Predictor:

- Concerns : i) Predict if an ins? is BR or not.
ii) if BR, predict its dir. (forward / backward)
iii) predict its target (EA, if taken)



① given an insⁿ. with some PC value,
 → it remains as BR or non-BR, throughout program runtime

↓
 + PC we remember if its a BR & its type:
 {
 → Unconditional BR TAKEN B
 → Conditional BR (need to Predict) BEQ, BNE
 (depending on result of a prev. CMP insⁿ, some flags are SET)
 → fn. call TAKEN call
 → return TAKEN iret
 → Non-BR

using a str. in H/w Insⁿ. Status table (IST)

provides better RANDOMNESS in case of Spatial & temporal locality

uses LSB 10-bits of PC, to mark the insⁿ as to which category it belongs.

dis-adv.: Destructive interference.

many addresses maps to same entry

keep updating the entries in IST.

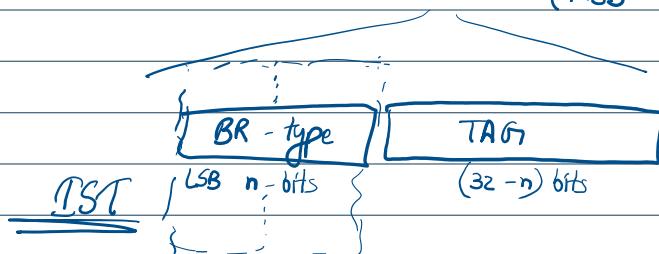
i.e., 1 PC being BR as lsb 10-bits, another being non-BR with same IST entry

BRANCH ALIASING

(due to small-sized IST)

Solⁿ: To distinguish, augment each entry of IST with the TAG value.

(MSB 32-n bits)



→ If for a PC, we don't have entry in IST,

we'll have to predict if its a BR or not.

→ LST works because of temporal locality exhibited by codes (i.e., LOOPS)

② Prediction :

→ Branches in loops have roughly similar behaviour most of the time.
Predominantly NOT TAKEN.

I] Simple BIMODAL Predictor: (1-bit Predictor)

+ PC in LST, it remembers the last outcome of conditional BR.

↳ 1, uses it for current Prediction.

Same augmented LST table,
along with 1-bit for BR as TAKEN / NOT TAKEN.
for last recorded outcome of BR.

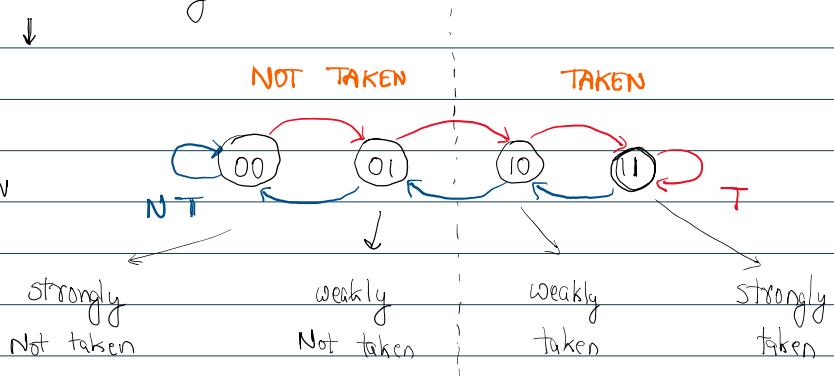
II] Saturating counter (2-bits predictor)

→ instead of having 1-bit history, uses 2-bits for hysteresis.

Algo. for update:

→ if a BR is TAKEN, ↑ saturating ctr.
else, ↓

for Prediction: 00, 01 → TAKEN
10, 11 → NOT TAKEN



→ works well when BR is biased, towards 1 dir.

{either mostly TAKEN or NT}

* Global history Register (GHR):

Consider, a shift-register, that records history of last n-branches encountered by processor & regardless of PC value {

1 bit for each BR ($0 \rightarrow \text{NT}$, $1 \rightarrow \text{Taken}$)

Hy,

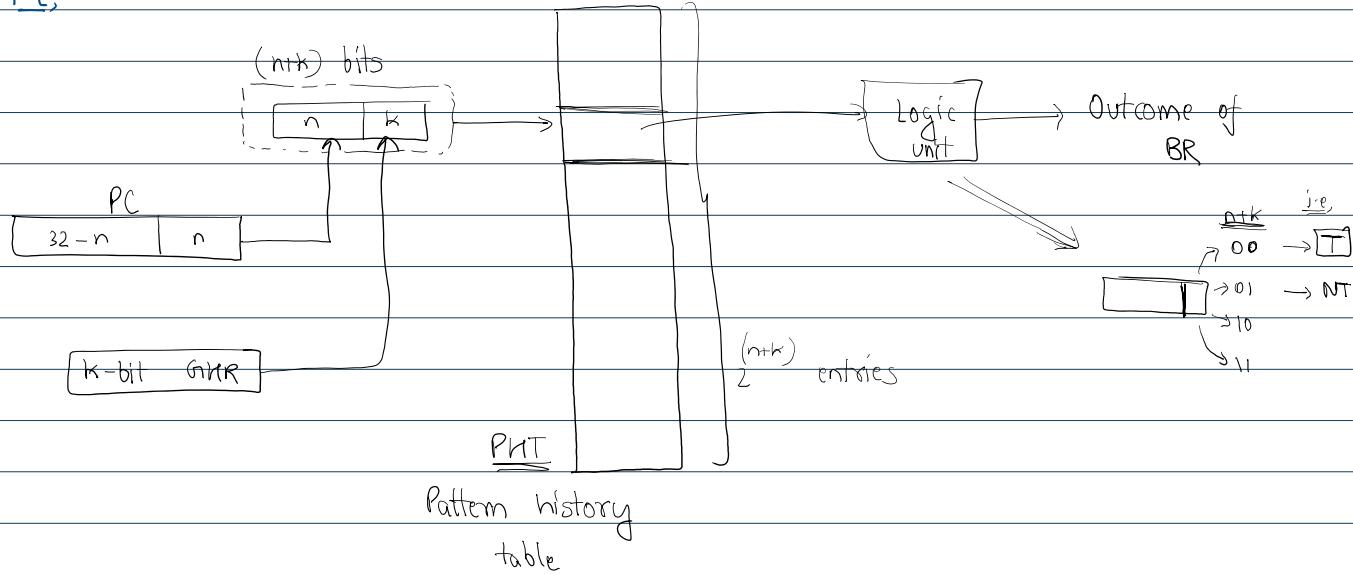
we have 2-bit shift-register (GMR)

↳ uses the history of last 2 BR ins?

III] GAp Predictor :

→ Use k-bit GMR along with Pattern History Table (PHT) to decide the outcome of a BR ins?

i.e.,



GMR : captures the behaviour of related branches.

i.e., {within a fn. scope}

IV] PAp Predictor :

→ uses multiple GMRs vs n-bits of PC,

→ Most Accurate.

to decide outcome of the BR.

V] Tournament Predictor :

very biased BR → bimodal predictor with saturating ctr., works well.

Alternating pattern → PAp predictor, works well

for a general piece of code,
we'll use 2 Predictors, I choose 1 of them
{based on behaviours}

→ To identify the target of BR insⁿ,

we augment the IST table \Rightarrow Branch Target Buffer (BTB)
i.e,

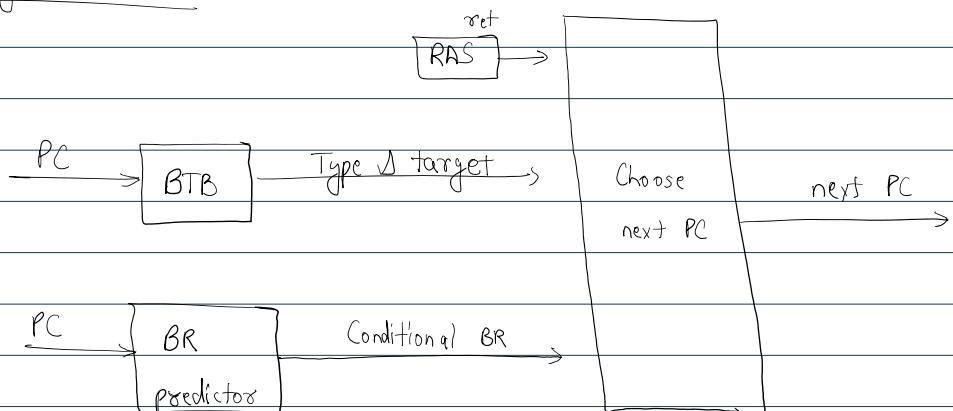
it now hold EA aswell.

Entry \rightarrow BR type | Tag-bits | Target or EA

→ for CALL / RETURN BR insⁿ,

we use Return Addr. Stack (RAS)

Summary of BR Pred.:



* Decode stage: (last part of IF unit)

Process of decoding:

- Expand imm. val. to 32/64 bit values.
- Extract all fields (IDs of regs.)
- Compute BR EA (if TAKEN) $\xrightarrow{\text{we usually have PC-relative branches}}$
- Add all implicit sources (i.e., ret. insⁿ)
- Create Insⁿ pkt. (includes Control-bits) $\xrightarrow{\text{return addr. [RA] reg.}}$

CISC insⁿ

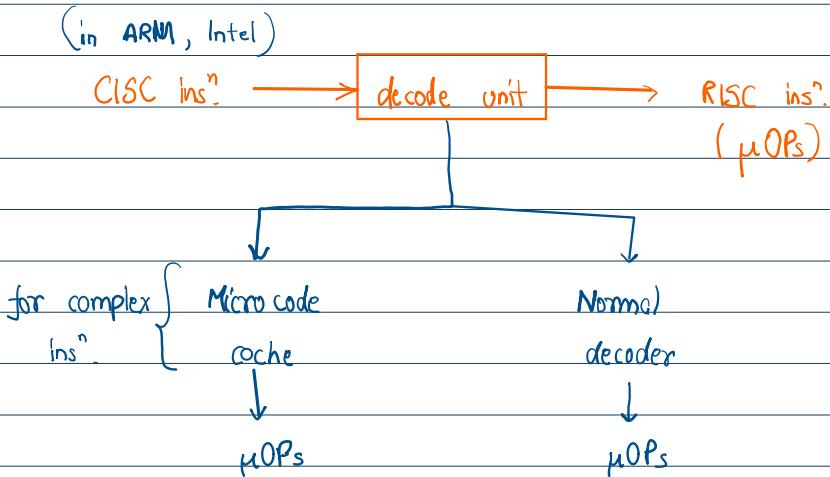
→ do not have fixed len.

i.e., in x86: len. varies from 1 to 15 Bytes.

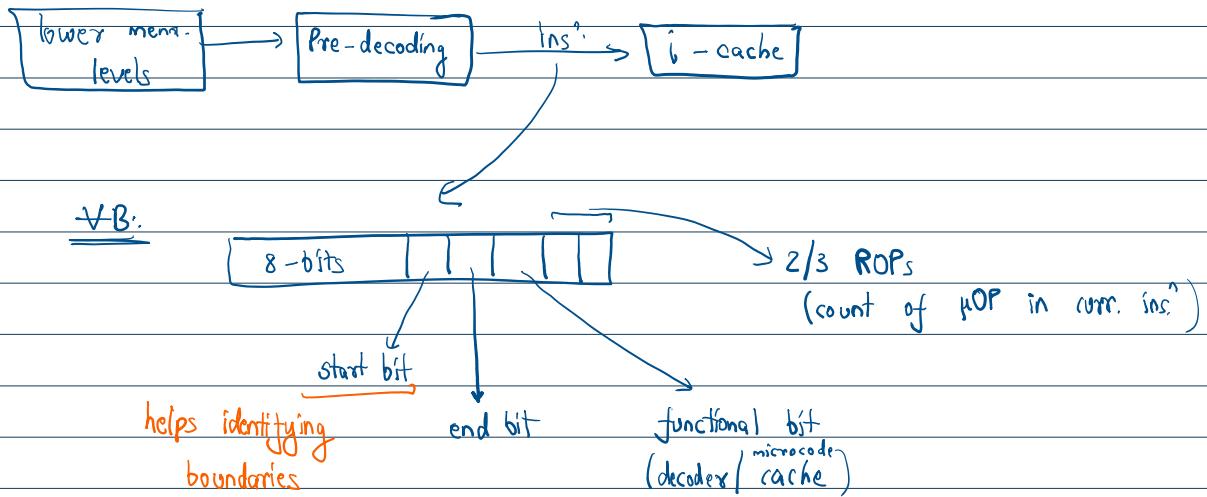
↓ SEQUENTIAL
hard to fetch multiple insⁿ at once, (\because linear scanning reqd.)

hard for OOO pipeline to process them.

hard for OOO pipeline to process them.
 CISC processors internally convert it to RISCy



→ To identify START/END of + CISC ins.,
 + Byte is augmented with some additional info.
 (done by L2 cache typically)



* ins. compression in i-cache:

- using shorter width ins. (with smaller ISA)
- replacing frequently executed ins. with codewords (using dictionary)

stores decoded ins.

Issue, Execute, Commit Stages in OOO pipeline:

Reg. renaming with Phy. reg. file: (done by H/w)

	#architectural reg. (visible to s/w)	
(intel, AMD) x86	8	→ used by compiler
x86-64	16	in ASSEMBLY CODE gen?
ARM	16	
MIPS	32	

MIPS: MicroProcessor w/o Interlocked Pipeline Stages.
 { INORDER pipeline }

more the #Phy. reg. → ↑ ILP can be achieved.
 (store interim values)

* Renaming uses 3 h/w structures:

i) Reg. Alias Table (RAT)

translates arch. reg. ID to phy. reg. ID.

ii) Dependency Check logic (DCL)

take care of dependencies b/w multiple ins'. renamed in same cycle,
 { using MUX }

iii) Free list/queue (circular Q)

Maintain a list of unmapped Phy. registers.

→ RAT also store avl. bit,

{ whether Phy. reg. value ready or not }

↳ if not, ins' has to WART in ins' window.

Rename table entry : Phy. reg. ID | avlbl. bit

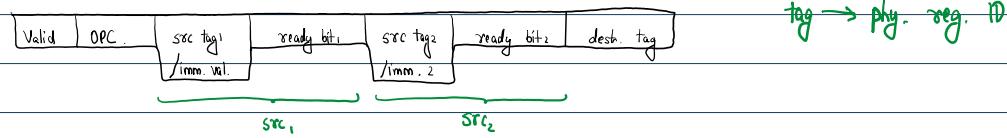
Dispatch:

(or issue) till rename stage, ins' proceeds INORDER

Sends the ins' to ins' window

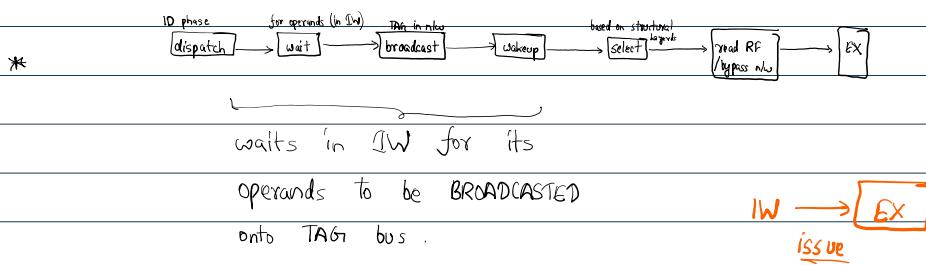


→ structure of a ins^n in fns^n window:



When both the operands are available, the ins^n is ready
to EXECUTED

either from Reg. file
/ Bypass n/w.

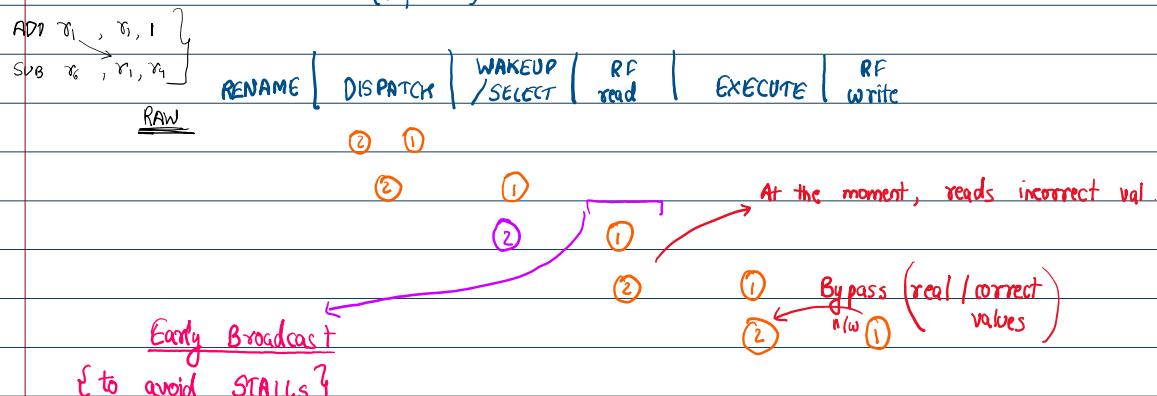


Once an ins^n has all its operands READY,
it asserts its attached REQUEST line

The Select unit chooses among these ready ins^n .
(GRANTS it)

→ Speculative Broadcasting:

Consider 2 ins^n with RAW dependency.
($i^{th}, i+1^{th}$)



Allows back-to-back EXⁿ of dependent ins^n

→ In-case of a 2 cycle LOAD-USE hazard:

where LOAD requires 2 cycles in EX phase,
we BROADCAST TAG 1 cycle later.

i.e., STALL of 1 cycle persists in ooo

i.e., for diff. insⁿ,
we BROADCAST TAG at diff. pt. of time.

→ During speculative BROADCAST, (Rename table)
we also send a signal to RAT table,
to update the avlbl. bit of processed reg. [Destn. reg]
{ in-order to help insⁿ that arrive after Broadcast, to not wait indefinitely }

* Problem:

→ insⁿ being written to the LW will miss the broadcast.
→ These insⁿ would also read avlbl. bit as 0, from RAT,
leading to indefinite WAITING

Solⁿ: i) Split 1st half of clk cycle for DISPATCH.
ii) use 2nd half for BROADCAST. } Not feasible

ii) Double BROADCAST

(w next cycle update)

→ log the BROADCASTED tags in a temporary str.
→ record all insⁿ that are being dispatched in a separate str.,
called **DISPATCH buffer**.

In the next cycle,

→ match the tag entries in the dispatch buffer,
→ if matched, update corresponding LW entry.
→ Clear entries of last cycle from buffer.

* Reg. Read & Execute:

LSQ : Load-Store Queue

↳ Enforces mem. dependencies.

Consider,

ld r₁, 4[r₃]
st r₂, 10[r₅] → if same MA.
↓
Problem

They're then sent to load-store unit:

- loads can execute immediately.
- stores update the processor's state (for Precise Exceptions)
 - { MUST be performed in Program Order }
 - ↪ can wait.

STORE Queue ≈ data-cache str.

- ↪ LOADS cannot be sent to mem. sys., until all the STORES before it are RESOLVED,
- I do not write to the same addr

* LSD:

<u>Op. type</u>	<u>Search dir.</u>	<u>Condition : Action</u>
LOAD	search all stores before it	<ul style="list-style-type: none"> → found store to same addr : terminate & forward values → found store with unresolved addr. : terminate → Both not met : goto d-cache.
STORE	search all LD/ST after it.	<ul style="list-style-type: none"> → found store with same addr. : terminate → found store with unresolved addr. : terminate → LOAD from same addr : forward value / addr.

* Ins? commit:

Reorder Buffer: (inorder Q str.)

- Contains 1 entry + insⁿ. fetched. (in program-order)
- After decoding the insⁿ, we enter it to ROB.
- if no free entry → ooo pipeline STALLS.

Entry is cleared when an insⁿ. COMMITs / Retires. → also, add reclaimed phy. reg. back to free list. (if any)

+ cycle, we inspect for top \boxed{W} entries, if READY to be commit.
in ROB

↪ Stop if any entry found to be NOT READY.

{ in general, fetch width = Commit width }

* Unexpected events :

- interrupts
- exceptions
- Branch mispredictions

↓

Some ins? in pipeline are on wrong path, i.e., they shouldn't be committed.

→ Mark ins? that has had an exception / suffered Misprediction in ROB.

In case of interrupt, mark topmost entry in ROB,

→ wait till marked ins? retires

→ initiate recovery seq. → flush ROB, LSQ, IW
↓ all pipeline reg..

+

Restore / Rollback arch. reg., state, PC

* VLIW Processor :

(Very long ins? word)

→ exploits ILP.

→ It relies on compile time analysis,
to identify ↓ bundle together ins? that can be executed
concurrently.

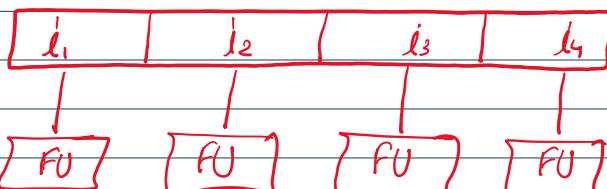
→ compiler is responsible for static scheduling of ins? in VLIW processors.

→ compiler finds out which operations can be executed in // in the program
↓, groups together these op? into single ins? known as ISSUE packet
or Bundle, which is a very large ins? word.

→ compiler ensures that an op? is NOT issued before its operands are ready.

→ They deploy multiple independent FU.
i.e.

Bundle :



(a VLIW ins?)

VLIW Processors

- Parallelism is performed STATICALLY at compile time, by Compiler.
- Data dependency is checked by compiler.
- less H/w.
- complex compiler.

Superscalar Processors

- dynamically at RUNTIME , by H/w
- data dependency is checked & resolved by H/w
- more H/w
- complex H/w